



# **SELF-CONSISTENCY IMPROVES CHAIN OF THOUGHT REASONING IN LANGUAGE MODELS**

Xuezhi Wang, Jason Wei , Dale Schuurmans, Quoc Le, Ed H. Chi,  
Sharan Narang, Aakanksha Chowdhery, Denny Zhou  
Google Research, Brain Team



# AGENDA

**1 MOTIVATION**

**2 RELATED WORKS**

**3 SELF-CONSISTENCY METHOD**

**4 EVALUATION**

**5 ADDITIONAL STUDIES**

**6 CONCLUSION**

**7 APPENDIX**



## CHAPTER 1

# MOTIVATION

# CHAIN-OF-THOUGHT REASONING

**CoT:** a series of prompts which mimic human reasoning to guide language models in their reasoning process.

**Q:** If there are 3 cars in the parking lot and 2 more cars arrive, how many cars are in the parking lot?”,

**A:** 5 

**CoT Reasoning:**

- There are 3 cars in the parking lot already.
- 2 more arrive.
- Now there are  $3 + 2 = 5$

The Answer is 5 

# COT REASONING

Scaling LMs and implementing CoT improves reasoning abilities for tackling complex tasks.

## Pros:

- CoT guides models in step-by-step reasoning, enhancing their ability to reason effectively.
- Achieve higher performance on complex tasks requiring multi-step reasoning.
- A transparent CoT, improves the interpretability of model decision.

## Possible to leverage CoT's benefits to achieve greater Consistency in finding the best solution? <sup>1</sup>

- **Consistency:** A desirable property of language understanding models.
- Improving overall language understanding and interpretation.
- Ensuring consistent performance in different linguistic situations.

<sup>1</sup> Measuring and Improving Consistency in Pretrained Language Models, [Elazar et al. \(2021\)](#)



## CHAPTER 2

# RELATED WORKS

## 1. Training Verifiers to Solve Math Word Problems, [Cobbe et al., 2021, Google](#)

### Challenge:

State-of-the-art language models struggle with multi-step mathematical reasoning.

### Idea:

Train an additional verifier to re-rank generated solutions.

- The paper introduces GSM8K, a dataset of diverse grade school math word problems.
- Sample a fixed number of candidate solutions, select the solution ranked highest by the verifier.
- Verifiers are trained to judge the correctness of model completions.
- Verification significantly improves performance on GSM8K.
- Improves the solve rate on math tasks compared to just fine-tuning the language model

## 1. Training Verifiers to Solve Math Word Problems, [Cobbe et al., 2021, Google](#) (continue)

### Cobbe et al., 2021

vs.

### Self-Consistency:

- Sample a fixed number of candidate solutions.
- Verifiers trained to judge the correctness of model completions
- Fine-tuning with human annotated reasoning paths
- Select the solution ranked highest by the verifier.



X

X



? Will be discussing



## **2. Measuring and Improving Consistency in Pretrained Language Models,** **Elazar et al. (2021)**

### **Challenge:**

Assess the consistency of Pretrained Language Models (PLMs) with respect to factual knowledge.

### **Idea:**

Create the PARAREL dataset to evaluate PLM consistency and propose a method for improving model consistency.

Enhance factual knowledge consistency through pre-training with additional consistency loss.

## 2. Measuring and Improving Consistency in Pretrained Language Models, [Elazar et al. 2021](#) (continue)

### Elazar et al. (2021)

vs.

### Self-Consistency

- Recognize the importance of consistency in LMs.
- Acknowledge the limitations of current language models in terms of consistency
- Improve consistency through additional training and experimentation.
- Proposes extending pre-training with an additional consistency loss to improve model consistency





### CHAPTER 3

# SELF-CONSISTENCY METHOD WITHIN THE COT

# SELF-CONSISTENCY METHOD


## Spec:

- Unsupervised:
- Does not rely on human annotation, additional training or need to further modifications.
- Acting as a self-ensemble approach without the need for multiple separately trained models.
- Improved performance without introducing additional complexity or training requirements.

## How?

- Complex reasoning tasks have multiple valid paths.
- Self-consistency encourages diverse problem-solving approaches.
- Thoughtful analysis expand the range of reasoning paths.
- It challenges the notion of a single "right" solution.
- Considering various paths enhances flexibility in finding the best solution.

# SELF-CONSISTENCY METHOD



**Q:** *If there are 3 cars in the parking lot and 2 more cars arrive, how many cars are in the parking lot?*

**A:** There are 3 cars in the parking lot already. 2 more arrive. Now there are  $3 + 2 = 5$  cars. The answer is 5.

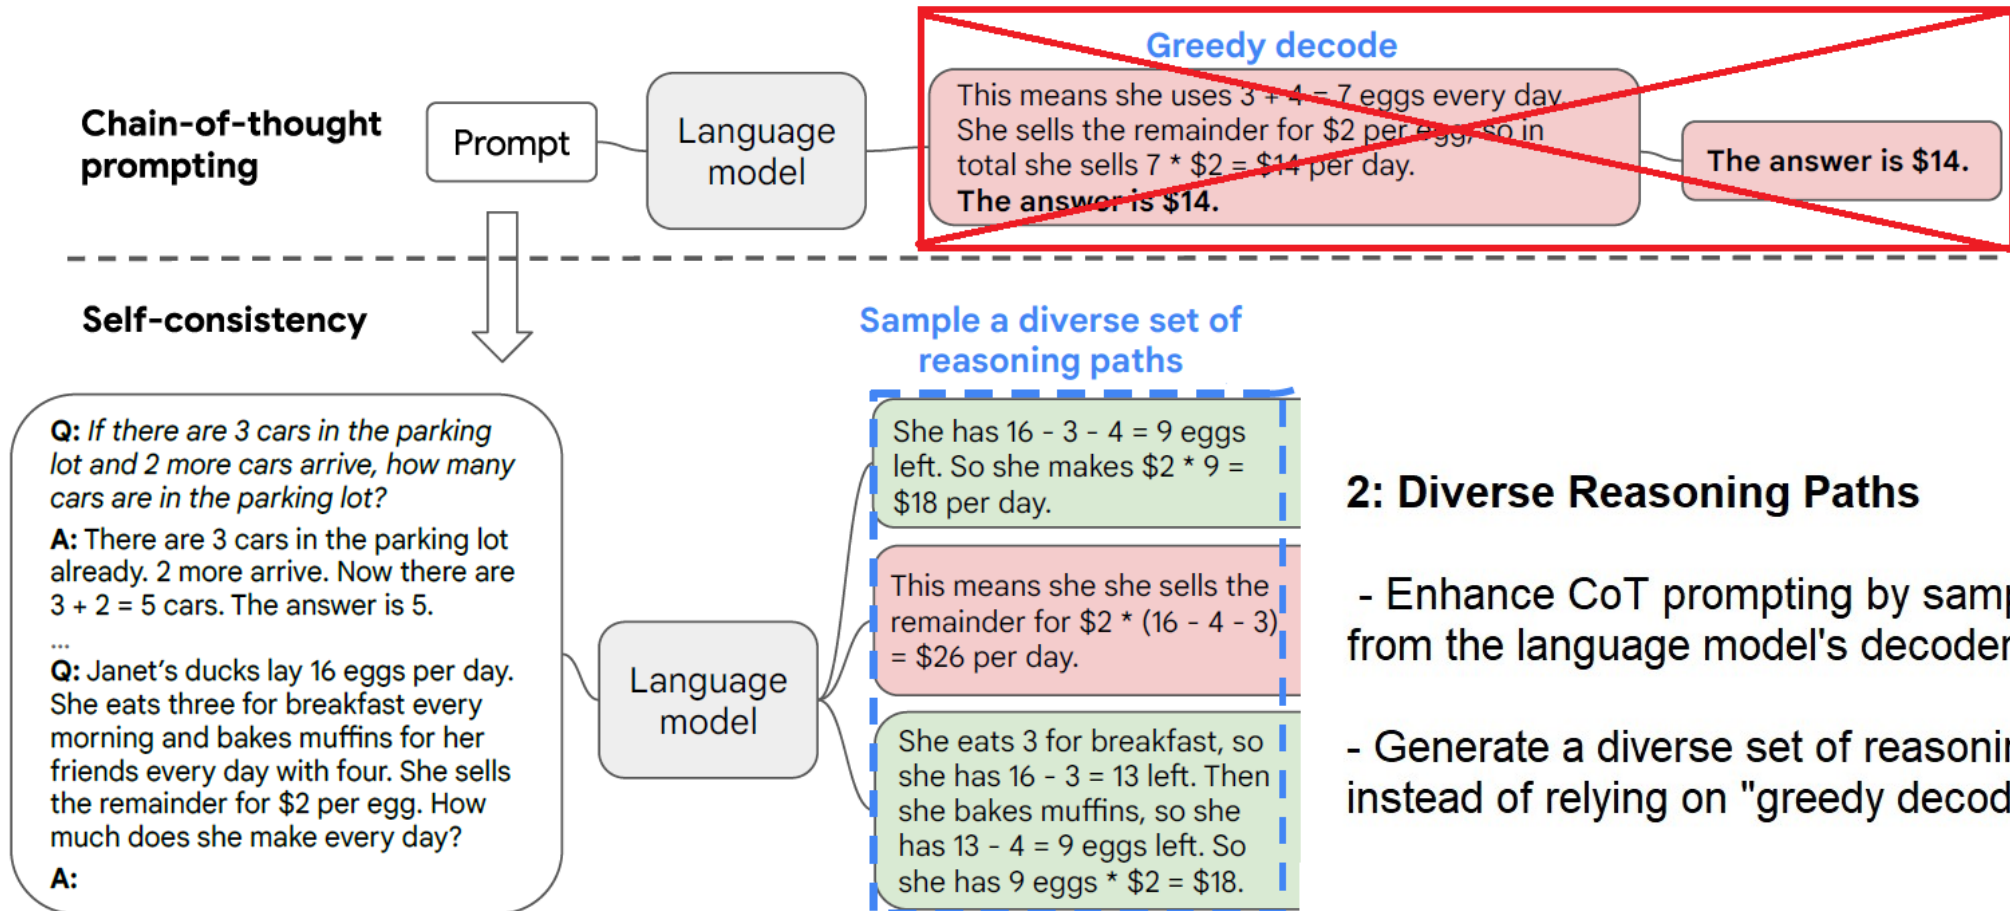
...

**Q:** Janet's ducks lay 16 eggs per day. She eats three for breakfast every morning and bakes muffins for her friends every day with four. She sells the remainder for \$2 per egg. How much does she make every day?

**A:**

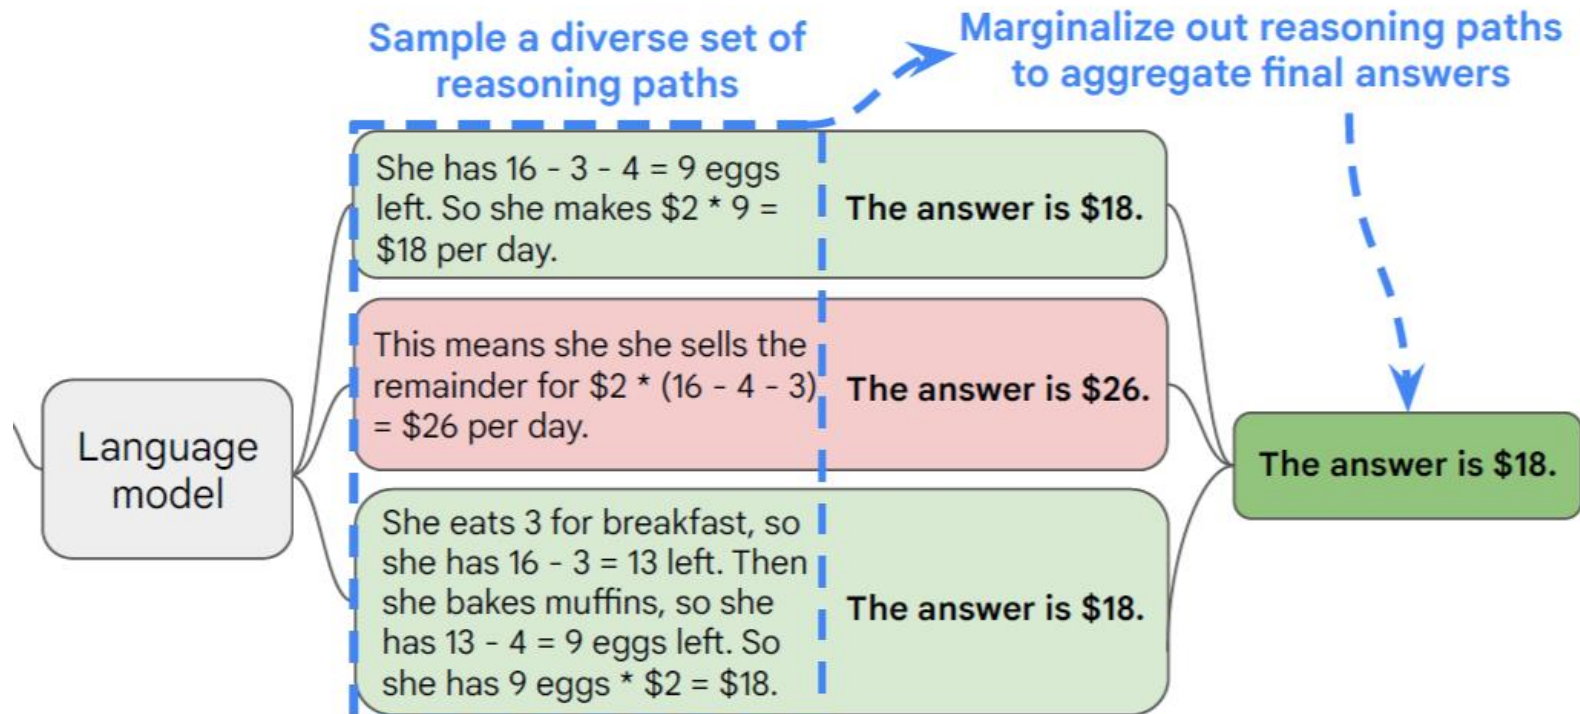
## 1: CoT Prompting

a language model is prompted with a set of manually written chain-of-thought exemplars



## 2: Diverse Reasoning Paths

- Enhance CoT prompting by sampling from the language model's decoder.
- Generate a diverse set of reasoning paths instead of relying on "greedy decode."



### 3. Aggregating Reasoning Paths for Consistent Answers

- Aggregate the reasoning paths by marginalizing them out.
- Select the most consistent answer from the final set of answers.

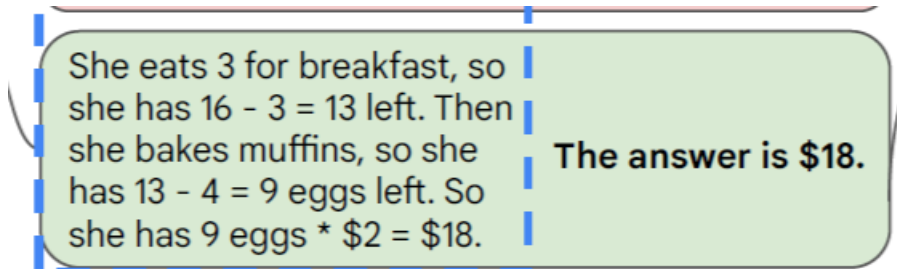
# SAMPLING

Self-consistency is compatible with most existing sampling algorithms, including:

- Temperature Sampling (Ackley et al., 1985; Ficler & Goldberg, 2017)
- Top-K Sampling (Fan et al., 2018; Holtzman et al., 2018; Radford et al., 2019)
- Nucleus Sampling (Holtzman et al., 2020)



# SELF-CONSISTENCY OVER DIVERSE REASONING PATHS



$r_i \rightarrow a_i$

- After Sampling multiple  $(r_i, a_i)$
- Apply marginalization over  $r_i$  by taking a **majority vote over  $a_i$**

$r_i$ : a sequence of tokens representing the reasoning path  
 $a_i$  the generated answers  $\in \mathbf{A}$ : fixed answer set  
 $m$ : # of candidate outputs sampled from the decoder

$$\arg \max_a \sum_{i=1}^m \mathbb{1}(a_i = a)$$

**Majority  
Vote**

**The most “consistent” answer  
among the final answer set.**

# Accuracy Comparison across commonsense reasoning benchmarks

	GSM8K	MultiArith	AQuA	SVAMP	CSQA	ARC-c
Greedy decode	56.5	94.7	35.8	79.0	79.0	85.2
Weighted avg (unnormalized)	$56.3 \pm 0.0$	$90.5 \pm 0.0$	$35.8 \pm 0.0$	$73.0 \pm 0.0$	$74.8 \pm 0.0$	$82.3 \pm 0.0$
Weighted avg (normalized)	$22.1 \pm 0.0$	$59.7 \pm 0.0$	$15.7 \pm 0.0$	$40.5 \pm 0.0$	$52.1 \pm 0.0$	$51.7 \pm 0.0$
Weighted sum (unnormalized)	$59.9 \pm 0.0$	$92.2 \pm 0.0$	$38.2 \pm 0.0$	$76.2 \pm 0.0$	$76.2 \pm 0.0$	$83.5 \pm 0.0$
Weighted sum (normalized)	$74.1 \pm 0.0$	$99.3 \pm 0.0$	$48.0 \pm 0.0$	$86.8 \pm 0.0$	$80.7 \pm 0.0$	$88.7 \pm 0.0$
Unweighted sum (majority vote)	$74.4 \pm 0.1$	$99.3 \pm 0.0$	$48.3 \pm 0.5$	$86.6 \pm 0.1$	$80.7 \pm 0.1$	$88.7 \pm 0.1$

Table 1: Accuracy comparison of different answer aggregation strategies on PaLM-540B.

# DIFFERENT ANSWER AGGREGATION STRATEGIES

## Weighted Aggregation

Weight each  $(r_i, a_i)$  by  $P(r_i, a_i \mid \text{prompt, question})$   
either take the:

- **unnormalized probability** of the model generating:  $P(r_i, a_i \mid \text{prompt, question})$

or

- **normalize the conditional probability** by the output length (Brown et al., 2020)

$$P(\mathbf{r}_i, \mathbf{a}_i \mid \text{prompt, question}) = \exp \frac{1}{K} \sum_{k=1}^K \log P(t_k \mid \text{prompt, question}, t_1, \dots, t_{k-1})$$

**Normalized  
weighted  
sum**

k: total # of tokens

log probability of generating the k-th token  $t_k$  in  $(r_i, a_i)$  conditioned on the previous tokens

# Accuracy Comparison across commonsense reasoning benchmarks

	GSM8K	MultiArith	AQuA	SVAMP	CSQA	ARC-c
Greedy decode	56.5	94.7	35.8	79.0	79.0	85.2
Weighted avg (unnormalized)	56.3 $\pm$ 0.0	90.5 $\pm$ 0.0	35.8 $\pm$ 0.0	73.0 $\pm$ 0.0	74.8 $\pm$ 0.0	82.3 $\pm$ 0.0
Weighted avg (normalized)	22.1 $\pm$ 0.0	59.7 $\pm$ 0.0	15.7 $\pm$ 0.0	40.5 $\pm$ 0.0	52.1 $\pm$ 0.0	51.7 $\pm$ 0.0
Weighted sum (unnormalized)	59.9 $\pm$ 0.0	92.2 $\pm$ 0.0	38.2 $\pm$ 0.0	76.2 $\pm$ 0.0	76.2 $\pm$ 0.0	83.5 $\pm$ 0.0
Weighted sum (normalized)	74.1 $\pm$ 0.0	99.3 $\pm$ 0.0	48.0 $\pm$ 0.0	86.8 $\pm$ 0.0	80.7 $\pm$ 0.0	88.7 $\pm$ 0.0
Unweighted sum (majority vote)	74.4 $\pm$ 0.1	99.3 $\pm$ 0.0	48.3 $\pm$ 0.5	86.6 $\pm$ 0.1	80.7 $\pm$ 0.1	88.7 $\pm$ 0.1

Table 1: Accuracy comparison of different answer aggregation strategies on PaLM-540B.



## CHAPTER 4

# EVALUATION

# EXPERMINET SETUP

## ▪ Benchmarks:

- **Arithmetic reasoning** (use the Math Word Problem Repository, ...)
- **Commonsense reasoning** (CommonsenseQA, StrategyQA, ...)
- **Symbolic Reasoning** (last letter concatenation, Coinflip )

- All experiments in the few-shot settings
- Neither training nor fine-tuning
- Use same prompts for fair comparison

## • Language Models:

UL2 (encoder-decoder, 20-B)<sup>1</sup>

LaMDA (decoder-only, 137-B)

GPT-3 (decoder-only, 175-B)<sup>2</sup>

PaLM (decoder-only, 540-B)

<sup>1</sup> [UL2](#)

<sup>2</sup> [GPT-3](#)

# MAIN RESULTS (ARITHMETIC REASONING)

	Method	AddSub	MultiArith	ASDiv	AQuA	SVAMP	GSM8K
	Previous SoTA	<b>94.9<sup>a</sup></b>	60.5 <sup>a</sup>	75.3 <sup>b</sup>	37.9 <sup>c</sup>	57.4 <sup>d</sup>	35 <sup>e</sup> / 55 <sup>g</sup>
UL2-20B	CoT-prompting	18.2	10.7	16.9	23.6	12.6	4.1
	Self-consistency	24.8 (+6.6)	15.0 (+4.3)	21.5 (+4.6)	26.9 (+3.3)	19.4 (+6.8)	7.3 (+3.2)
LaMDA-137B	CoT-prompting	52.9	51.8	49.0	17.7	38.9	17.1
	Self-consistency	63.5 (+10.6)	75.7 (+23.9)	58.2 (+9.2)	26.8 (+9.1)	53.3 (+14.4)	27.7 (+10.6)
PaLM-540B	CoT-prompting	91.9	94.7	74.0	35.8	79.0	56.5
	Self-consistency	93.7 (+1.8)	99.3 (+4.6)	81.9 (+7.9)	48.3 (+12.5)	86.6 (+7.6)	74.4 (+17.9)
GPT-3 Code-davinci-001	CoT-prompting	57.2	59.5	52.7	18.9	39.8	14.6
	Self-consistency	67.8 (+10.6)	82.7 (+23.2)	61.9 (+9.2)	25.6 (+6.7)	54.5 (+14.7)	23.4 (+8.8)
GPT-3 Code-davinci-002	CoT-prompting	89.4	96.2	80.1	39.8	75.8	60.1
	Self-consistency	91.6 (+2.2)	<b>100.0</b> (+3.8)	<b>87.8</b> (+7.6)	<b>52.0</b> (+12.2)	<b>86.8</b> (+11.0)	<b>78.0</b> (+17.9)

Table 2: Arithmetic reasoning accuracy by self-consistency compared to chain-of-thought prompting

# MAIN RESULTS (COMMONSENSE AND SYMBOLIC REASONING)

	Method	CSQA	StrategyQA	ARC-e	ARC-c	Letter (4)	Coinflip (4)
	Previous SoTA	<b>91.2<sup>a</sup></b>	73.9 <sup>b</sup>	86.4 <sup>c</sup>	75.0 <sup>c</sup>	N/A	N/A
UL2-20B	CoT-prompting	51.4	53.3	61.6	42.9	0.0	50.4
	Self-consistency	55.7 (+4.3)	54.9 (+1.6)	69.8 (+8.2)	49.5 (+6.8)	0.0 (+0.0)	50.5 (+0.1)
LaMDA-137B	CoT-prompting	57.9	65.4	75.3	55.1	8.2	72.4
	Self-consistency	63.1 (+5.2)	67.8 (+2.4)	79.3 (+4.0)	59.8 (+4.7)	8.2 (+0.0)	73.5 (+1.1)
PaLM-540B	CoT-prompting	79.0	75.3	95.3	85.2	65.8	88.2
	Self-consistency	80.7 (+1.7)	<b>81.6</b> (+6.3)	<b>96.4</b> (+1.1)	<b>88.7</b> (+3.5)	70.8 (+5.0)	91.2 (+3.0)
GPT-3 Code-davinci-001	CoT-prompting	46.6	56.7	63.1	43.1	7.8	71.4
	Self-consistency	54.9 (+8.3)	61.7 (+5.0)	72.1 (+9.0)	53.7 (+10.6)	10.0 (+2.2)	75.9 (+4.5)
GPT-3 Code-davinci-002	CoT-prompting	79.0	73.4	94.0	83.6	70.4	99.0
	Self-consistency	81.5 (+2.5)	79.8 (+6.4)	96.0 (+2.0)	87.5 (+3.9)	<b>73.4</b> (+3.0)	<b>99.5</b> (+0.5)





Table 3: Commonsense and symbolic reasoning accuracy by self-consistency compared to CoT prompting



# EFFECT OF THE NUMBER OF SAMPLED REASONING PATHS

Number of reasoning paths  
 $\approx$   
Accuracy

● Greedy Decode (Single-path)  
■ Self Consistency (Multi-path)





## CHAPTER 5

# ADDITIONAL STUDIES

# IMPROVING ROBUSTNESS TO IMPERFECT PROMPTS

- **Imperfect Prompts:** Manually constructed prompts in few-shot learning can contain minor mistakes due to human annotation.
- Greedy decoding with imperfect prompts leads to decreased accuracy (17.1% → 14.9%).
- Self-consistency fills in the gaps and significantly improves results with imperfect prompts.

LaMDA-137B	Prompt with correct chain-of-thought	17.1
	Prompt with imperfect chain-of-thought	14.9
	+ Self-consistency (40 paths)	<b>23.4</b>

- **Consistency and Accuracy:** Consistency, measured as the agreement with the final aggregated answer, is highly correlated with accuracy.




Figure 5: The consistency is correlated with model's accuracy.

# SELF CONSISTENCY FOR NON-NATURAL-LANGUAGE REASONING AND ZERO-SHOT COT

- Self-consistency improves accuracy by generating intermediate equations

Gains are smaller due to limited diversity in equation decoding.

e.g., from “There are 3 cars in the parking lot already. 2 more arrive.

Now there are  $3 + 2 = 5$  cars.” to “ $3 + 2 = 5$ ”) (**Limited Leeway**)

LaMDA-137B	Prompt with equations	5.0
	+ Self-consistency (40 paths)	<b>6.5</b>
PaLM-540B	Zero-shot CoT (Kojima et al., 2022)	43.0
	+ Self-consistency (40 paths)	<b>69.2</b>

- Self-consistency enhances results significantly in zero-shot CoT scenarios (+26.2%)



## CHAPTER 6

# CONCLUSION

# Conclusion

- Self-Consistency improves accuracy in a range of Arithmetic and Commonsense reasoning tasks across four language models.
- With reasoning paths enhances interpretability in reasoning tasks.
- Provides improved output calibration and uncertainty estimation.
- Computation cost: Self-Consistency incurs additional computational overhead.
- Optimal paths: Few reasoning paths yield significant gains without excessive cost.

# Future Work

- Use self-consistency to generate better supervised data for fine-tuning.
- Improving prediction accuracy in a single inference run.
- Mitigate inconsistencies and inaccuracies in reasoning paths to enhance trustworthiness.
- Further research needed to refine the process of generating effective rationales.

THANKS



# THANKS!



Ali Khalaji



ali.khalaji@stud.tu-darmstadt.de



**CHAPTER 7**

**APPENDIX**

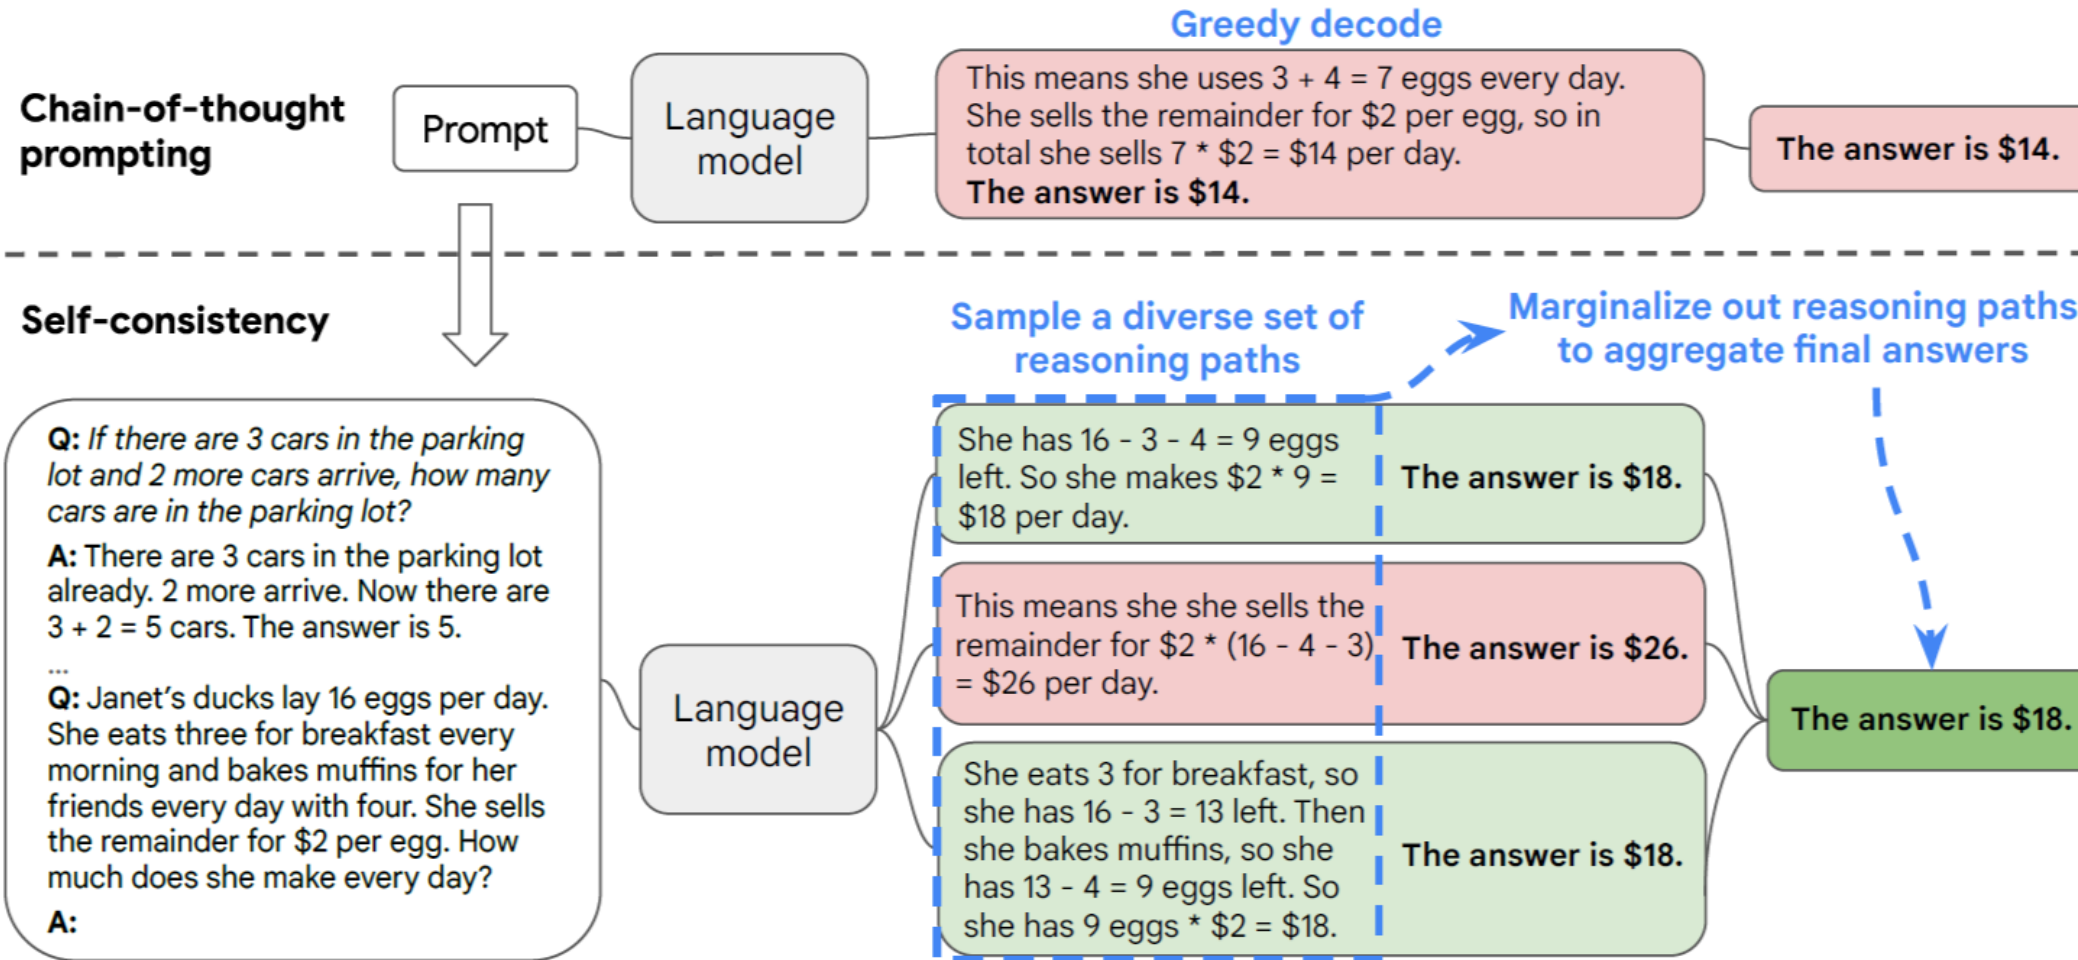


# CONSISTENCY IN LANGUAGE UNDERSTANDING <sup>1</sup>

- Possible to leverage CoT' s benefits to achieve greater Consistency in finding the best solution?
- **Consistency:** A desirable property of language understanding models.
- We want to make consistent decisions in semantically equivalent contexts.
- Improving overall language understanding and interpretation.
- Ensuring consistent performance in different linguistic situations.
- Promoting reliable and consistent language processing.

<sup>1</sup> Measuring and Improving Consistency in Pretrained Language Models, [Elazar et al. \(2021\)](#)

## APPENDIX - SELF-CONSISTENCY WITHIN THE COT REASONING



### Additional Notes on Self-Consistency

- Find a middle ground between open-ended and fixed answer text generation.
- Reasoning Tasks: Usually rely on greedy decoding approaches due to fixed answers.
- Benefits of Diversity: Adding diversity to reasoning processes highly advantageous.
- Sampling Approach: like in open-ended text generation, to introduce diversity.
- Self-consistency can be expanded to open-text generation if we define a suitable consistency metric.
- Metric Definition: Develop a metric to assess agreement or contradiction between multiple generated texts.

# **SELF-CONSISTENCY VS. SAMPLE-AND-RANK**

## **Experiment Setup**

- Comparison conducted on GPT-3 code-davinci-001 model.
- Same number of sequences sampled from the decoder for both approaches.
- Final answer extracted from the top-ranked sequence.

## **Self-Consistency Method**

- Generates multiple sequences from the decoder.
- Promotes agreement and consistency among the generated responses.
- Results in significant accuracy improvement.

## **Sample-and-Rank Method<sup>1</sup>**




- Sampling multiple sequences and ranking them based on log probability.
- Provides a marginal accuracy improvement.
- Gain is comparatively smaller compared to self-consistency.

<sup>1</sup> [Towards a Human-like Open-Domain Chatbot, \(Adiwardana et al., 2020\)](#)

## SELF-CONSISTENCY VS. SAMPLE-AND-RANK

### Results:

- Self-consistency proves to be a more effective approach for improving generation quality compared to the sample-and-rank method.
- Self-consistency yields substantial accuracy gains by promoting diversity and consistency among generated responses.



## SELF-CONSISTENCY VS. BEAM SEARCH

- on the UL2-20B Model.
- Equal number of beams (for beam search) and reasoning paths (for self-consistency)
- Self-consistency can adopt beam search to decode each reasoning path
- The Performance of Self-Consistency with sampling is better, as beam search's limited output diversity limits its effectiveness.<sup>1</sup>

Beam size / Self-consistency paths		1	5	10	20	40
AQuA	Beam search decoding (top beam)	23.6	19.3	16.1	15.0	10.2
	Self-consistency using beam search	23.6	19.8 $\pm$ 0.3	21.2 $\pm$ 0.7	24.6 $\pm$ 0.4	24.2 $\pm$ 0.5
	Self-consistency using sampling	19.7 $\pm$ 2.5	<b>24.9 <math>\pm</math> 2.6</b>	<b>25.3 <math>\pm</math> 1.8</b>	<b>26.7 <math>\pm</math> 1.0</b>	<b>26.9 <math>\pm</math> 0.5</b>
MultiArith	Beam search decoding (top beam)	10.7	12.0	11.3	11.0	10.5
	Self-consistency using beam search	10.7	11.8 $\pm$ 0.0	11.4 $\pm$ 0.1	12.3 $\pm$ 0.1	10.8 $\pm$ 0.1
	Self-consistency using sampling	9.5 $\pm$ 1.2	11.3 $\pm$ 1.2	<b>12.3 <math>\pm</math> 0.8</b>	<b>13.7 <math>\pm</math> 0.9</b>	<b>14.7 <math>\pm</math> 0.3</b>

Compare self-consistency with beam search decoding on the UL2-20B model.

<sup>1</sup> [A Simple, Fast Diverse Decoding Algorithm for Neural Generation, Li & Jurafsky, 2016\)](#)

## SELF CONSISTENCY VS. CHAIN OF THOUGHT

- For some tasks (e.g., ANLI-R1, e-SNLI, RTE), adding chain-of-thought does hurt performance compared to standard prompting, but
- but self-consistency is able to robustly boost the performance and outperform standard prompting, making it a reliable way to add rationales in few-shot in-context learning for common NLP tasks

	ANLI R1 / R2 / R3	e-SNLI	RTE	BoolQ	HotpotQA (EM/F1)
Standard-prompting (no-rationale)	69.1 / 55.8 / 55.8	85.8	84.8	71.3	27.1 / 36.8
CoT-prompting (Wei et al., 2022)	68.8 / 58.9 / 60.6	81.0	79.1	74.2	28.9 / 39.8
Self-consistency	<b>78.5 / 64.5 / 63.4</b>	<b>88.4</b>	<b>86.3</b>	<b>78.4</b>	<b>33.8 / 44.6</b>

Table 5: Compare Standard/CoT prompting with self-consistency on common NLP tasks.

**THE END**

---

**THE END**